

MA589 Project 1

Megha Pandit

September 26, 2018

1. (WaRming up) Write (R) functions that return:

- (a) The inverse or the transpose inverse of an upper triangular matrix. Call this function `inv.upper.tri` and provide a transpose argument to specify if the transpose is requested. Hint: use `backsolve`.

```
inv.upper.tri<-function(A = matrix, v = matrix, transpose){  
  backsolve(r = A,x = v,transpose = transpose)  
}
```

(b) Quick check: if `u <- 1e200 * rep(1, 100)`, what is `norm2(u)`?

If we directly do a `crossprod()` on `u`, we get `Inf` as the answer, since `1e200` is a large value. Therefore, we can multiply and divide the each element of the vector `u` by the maximum value in the vector and then find the `norm2` of `u`.

```
u <- 1e200*rep(1,100)  
norm2 <- function(v){  
  max_v <- max(abs(v))  
  sqrt(crossprod(v/max(v), v/max(v)))*max(v)  
}  
norm2(u)
```

```
##      [,1]  
## [1,] 1e+201
```

- (c) The column-normalization `U` of matrix `A`, $U_{ij} = A_{ij}/||A_j||$ (call this function `normalize.cols`, and feel free to use `norm2` above).

```
A <- matrix(1:20, 5, 4)  
normalize.cols <- function(A){  
  U <- matrix(0, nrow = dim(A)[1], ncol = dim(A)[2])  
  for (i in 1:dim(A)[1]){  
    for (j in 1:dim(A)[2]) {  
      U[i,j] <- A[i,j]/norm2(A[,j]) #using the norm2 function from part (b)  
    }  
  }  
  U  
}  
normalize.cols(A)
```

```
##      [,1]      [,2]      [,3]      [,4]  
## [1,] 0.1348400 0.3302891 0.3761921 0.3963019  
## [2,] 0.2696799 0.3853373 0.4103913 0.4210708  
## [3,] 0.4045199 0.4403855 0.4445906 0.4458397  
## [4,] 0.5393599 0.4954337 0.4787899 0.4706085  
## [5,] 0.6741999 0.5504819 0.5129892 0.4953774
```

- (d) Quick check: what is `proj(1:100, u)`, `u` as in (b) above?

For very large values $\text{proj}_u(a) = \frac{u^T a}{\|u\|^2} u$ tends to zero as $\|u\|^2$ tends to infinity. Therefore, we can split the expression as $\text{proj}_u(a) = \frac{u^T a}{\|u\|} \frac{u}{\|u\|}$.

```

a <- 1:100
u <- 1e200 * rep(1, 100)
proj_1 <- function(a, u) {
  (drop(crossprod(u, a)) / drop(norm2(u))) * (u / drop(norm2(u)))
}
proj_1(a = a, u = u)

## [1] 50.5 50.5 50.5 50.5 50.5 50.5 50.5 50.5 50.5 50.5 50.5 50.5 50.5 50.5
## [15] 50.5 50.5 50.5 50.5 50.5 50.5 50.5 50.5 50.5 50.5 50.5 50.5 50.5 50.5
## [29] 50.5 50.5 50.5 50.5 50.5 50.5 50.5 50.5 50.5 50.5 50.5 50.5 50.5 50.5
## [43] 50.5 50.5 50.5 50.5 50.5 50.5 50.5 50.5 50.5 50.5 50.5 50.5 50.5 50.5
## [57] 50.5 50.5 50.5 50.5 50.5 50.5 50.5 50.5 50.5 50.5 50.5 50.5 50.5 50.5
## [71] 50.5 50.5 50.5 50.5 50.5 50.5 50.5 50.5 50.5 50.5 50.5 50.5 50.5 50.5
## [85] 50.5 50.5 50.5 50.5 50.5 50.5 50.5 50.5 50.5 50.5 50.5 50.5 50.5 50.5
## [99] 50.5 50.5

```

(e) The Vandermonde matrix of vector $a = [a_i]_{i=1,\dots,n}$ and degree d

```

vander_monde <- function(a,d){
  V <- matrix(0, nrow = length(a), ncol = d+1)
  for (j in 1:(d+1)){
    for (i in 1:length(a)){
      V[i,j] <- a[i]^(j-1)
    }
  }
  V
}
c <- rep(1:4)
vander_monde(c,4)

```

```

##      [,1] [,2] [,3] [,4] [,5]
## [1,]    1    1    1    1    1
## [2,]    1    2    4    8   16
## [3,]    1    3    9   27   81
## [4,]    1    4   16   64  256

```

2. The machine epsilon, ϵ , can be defined as the smallest floating point (with base 2) such that $1 + \epsilon > 1$, that is, $1 + \epsilon/2 == 1$ in machine precision.

(a) Write a function that returns this value by starting at $\text{eps} = 1$ and iteratively dividing by 2 until the definition is satisfied.

```

machine_ep <- function(eps){
  while (1 + (eps/2) != 1) {
    eps <- eps/2
  }
  eps
}
machine_ep(1)

```

```
## [1] 2.220446e-16
```

(b) Write a function that computes $f(x) = \log(1 + \exp(x))$ and then evaluate: $f(0)$, $f(???)$, $f(80)$, and $f(800)$.

```

fun_x <- function(x){
  log(1 + exp(x))
}

```

```
}  
fun_x(0)
```

```
## [1] 0.6931472
```

```
fun_x(-80)
```

```
## [1] 0
```

```
fun_x(80)
```

```
## [1] 80
```

```
fun_x(800)
```

```
## [1] Inf
```

Since $\exp(-80)$ is a very small number almost equal to 0, $\log(1+0) = \log(1) = 0$. Hence, $\text{fun_x}(-80)$ is zero. Similarly, since $\exp(800)$ is a very large number tending to infinity, $\text{fun_x}(800)$ yields Inf as the answer.

- (c) How would you specify your function to avoid computations if $x \ll 0$ ($x < 0$ and $|x|$ is large)? (Hint: ϵ .)

Since machine epsilon is defined as the smallest floating point, for any $x < \text{epsilon}/2$, we can have $f(x)=0$

```
fun_small <- function(x1){  
  if(exp(x1) > machine_ep(x1)/2){  
    log(1 + exp(x1))  
  }  
  else{  
    log(1)  
  }  
}  
fun_small(-80)
```

```
## [1] 0
```

- (d) How would you implement your function to not overflow if $x \gg 0$?

To avoid computations when $x \gg 0$, we consider the smallest value of x for which $\log(1 + \exp(x)) = \log(\exp(x))$. We define s be such that $\log(1 + \exp(s)) = \log(\exp(s))$. If $x > s$, then we just return x without calculating $f(x)$.

```
fun_large <- function(x){  
  s<- 1  
  while(exp(s) != exp(s)+1){  
    s = s+1  
  }  
  s  
  if(x < s){  
    fun_x(x)  
  }else{  
    x  
  }  
}  
fun_large(800)
```

```
## [1] 800
```

3.(a) Show that $C = Q^T A$ is upper triangular and that C is the Cholesky factor of $A^T A$.

If A is a positive definite matrix, then we can find an upper triangular matrix C such that $A = C^T C$. This process is called Cholesky Decomposition. Applying QR Decomposition on A , we have $A = QR$ where Q is an orthogonal matrix and R is an upper triangular matrix. Given $C = Q^T A$, Q is an orthogonal matrix and hence $Q^{-1} = Q^T$.

$$C = Q^T A = Q^{-1} A$$

$$C = Q^{-1} QR = R$$

Since R is an upper triangular matrix and $C = R$, C is an upper triangular matrix. To show that C is the Cholesky factor of $A^T A$, we need to show that $A^T A = C^T C$. We have $C = Q^T A$.

$$C^T C = (Q^T A)^T (Q^T A)$$

$$C^T C = A^T Q Q^T A = A^T Q Q^{-1} A = A^T I A = A^T A$$

Therefore, $A^T A = C^T C$ and C is the Cholesky factor of $A^T A$.

3.(b) Write a R function that computes the Q orthogonal factor of a Vandermonde matrix with base vector x and degree d without computing the Vandermonde matrix explicitly, that is, as your function iterates to compute u_i , compute and use the columns of the Vandermonde matrix on the fly.

```
q_fun <- function(a,d){
  U <- matrix(nrow = length(a), ncol = d+1)
  e <- rep(1,d+1)
  x <- rep(0,d+1)
  U[,1] <- rep(1, length(a))
  e[2] <- crossprod(U[,1],U[,1])
  for (i in 2:(d+1)){
    sum_proj <- 0
    U[,i] <- a^(i-1)
    for (j in 1:(i-1)){
      proj_u <- proj_1(U[,i],U[,j])
      sum_proj <- sum_proj + proj_u
    }
    U[,i] <- U[,i] - sum_proj
    e[i+1] <- drop(crossprod(U[,i],U[,i])) #Calculating eta for part (c)
    x[i] <- t(U[,i])%%diag(a) %%U[,i]/(drop(crossprod(U[,i],U[,i]))) # Calculating alpha for part(c)
  }
  Q <- normalize.cols(U)
  return (list(Q,e,x))
}
a <- c(1,2,3)
d <- 3
Q_ortho <- q_fun(a,d)
Q_ortho
```

```
## [[1]]
##           [,1]           [,2]           [,3]           [,4]
## [1,] 0.5773503 -7.071068e-01  0.4082483 -0.3165797
## [2,] 0.5773503 -3.140185e-16 -0.8164966  0.8020019
## [3,] 0.5773503  7.071068e-01  0.4082483 -0.5065275
##
## [[2]]
## [1] 1.000000e+00 3.000000e+00 2.000000e+00 6.666667e-01 7.083971e-27
##
```

```
## [[3]]
## [1] 0.000000 2.000000 2.000000 2.156347
```

#For part(c) of the question

```
q <- Q_ortho[[1]]
e <- Q_ortho[[2]]
x <- Q_ortho[[3]]
```

3.(c) Write a R function that, given ?? and ??, computes Q.

```
x
```

```
## [1] 0.000000 2.000000 2.000000 2.156347
```

```
q_comp <- function ( e = vector , x = vector , a = vector ){
  Q <- matrix(0, nrow = length(a), ncol = d+1 )
  Q[,1] <- 1
  Q[,2] <- a - x[1]*rep(1,length(a))
  for ( i in 2:d){
    for ( j in 1:length(a)){
      Q[j,i+1] <- ((a[j] - x[i]) * Q[j,i]) - (e[i+1]/e[i] * Q[j,i-1]) #algorithm from the question
    }
  }
  return(Q)
}

Q <- q_comp(e,x,a)
Q_new <- normalize.cols(Q)
Q_new
```

```
##           [,1]      [,2]      [,3]      [,4]
## [1,] 0.5773503 0.2672612 -0.5661385 0.6666667
## [2,] 0.5773503 0.5345225 -0.2264554 -0.3333333
## [3,] 0.5773503 0.8017837 0.7925939 0.6666667
```

3.(d)

#Alpha1 is the mean of vector a

```
x[1] == mean(a)
```

```
## [1] FALSE
```

#Eta2 is the number of vectors in vector a

```
e[2] == length(a)
```

```
## [1] TRUE
```

*#Eta3 gives the value of (n-1)*Var(a)*

```
e[3] == (length(a) -1 )*var(a)
```

```
## [1] TRUE
```

4. (a) To prove that $H_0 : \beta_j = \beta_{j+1} = \dots = \beta_p = 0$ is equivalent to testing $\gamma_j = \dots = \gamma_p = 0$, where $\gamma = R\beta$

Since X has thin QR Decomposition, $X = QR$,

$$\beta = (X^T X)^{-1} X^T y = [(QR)^T QR]^{-1} (QR)^T y$$

$$\beta = [R^T Q^T QR]^{-1} R^T Q^T y$$

Since Q is an orthogonal matrix, $Q^T Q = I$

$$\beta = (R^T R)^{-1} R^T Q^T y = R^{-1} R^{-T} R^T Q^T y = R^{-1} Q^T y$$

Therefore,

$$R\beta = Q^T y = (Q^T Q)^{-1} Q^T y$$

Since $\gamma = R\beta$,

$$\gamma = Q^T y$$

, H_0 is equivalent to testing $\gamma_j = \dots = \gamma_p = 0$ Hence, $y = Q\gamma$ and y can be regressed on Q instead of X . And,

$$\text{Var}(\gamma) = \text{Var}(Q^T y) = E((Q^T y)(Q^T y)^T) = Q^T E(Y^T Y) Q = \sigma^2 I_n Q^T Q = \sigma^2 I_n$$

4.(b) Show that the ML estimator for γ is $\hat{\gamma} = Q^T y$ and the components of $\hat{\gamma}$ are independent.

Since $y \sim N(Q\gamma, \sigma^2 I_n)$, $y = \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{1}{2\sigma^2}(y-Q\gamma)^2}$ So, $y = (\text{constant})e^{-(y-Q\gamma)^2} = (\text{constant})e^{-(y-Q\gamma)^T(y-Q\gamma)}$

To maximize the expectation, we need to minimize $f = (y - Q\gamma)^T(y - Q\gamma)$ Equating $\frac{\partial f}{\partial \gamma}$ to zero, we get

$$\frac{\partial(y^T y - 2Q^T y \gamma + (Q\gamma)^T(Q\gamma))}{\partial \gamma} = 0, \text{ i.e., } -2Q^T y + 2\hat{\gamma} = 0$$

Therefore,

$$\hat{\gamma} = Q^T y$$

In order to prove that the components of $\hat{\gamma}$ are independent, we can show that the covariance or the non-diagonal terms of the $\text{Cov}[\hat{\gamma}]$ are zero.

$$\text{cov}[\hat{\gamma}] = \text{cov}[Q^T y] = E((Q^T y)(Q^T y)^T) = E[Q^T y y^T Q]$$

$$\text{cov}[\hat{\gamma}] = Q^T E(y y^T) Q = Q^T \sigma^2 I_n Q = \sigma^2 I_n$$

$\sigma^2 I_n$ matrix has its non-diagonal elements equal to zero, and hence, the components of $\hat{\gamma}$ are independent

4.(c) Using R, explain how you compute: (i) the ML estimate $\hat{\gamma}$ as a function of $\hat{\beta}$, and (ii) the correlation matrix of $\hat{\gamma}$ using only `crossprod`, `normalize.cols`, and `inv.upper.tri`.

Since $\gamma = R\beta$,

$$\hat{\beta} = R^{-1} \hat{\gamma}$$

In R, we can use the `inv.upper.tri` function that we defined in question 1. `#b = beta hat and g = gamma vector and R is the upper triangular matrix` `b <- inv.upper.tri(R, g, transpose = FALSE)`

In R,

$$\text{cor}(\hat{\beta}) = \text{crossprod}(\text{normalized.cols}(\hat{\beta}), \text{normalized.cols}(\hat{\beta}))$$

$$\text{cor}(\hat{\beta}) = \text{crossprod}(\text{normalized.cols}(\text{inv.upper.tri}(R, \hat{\gamma}), \text{normalized.cols}(\text{inv.upper.tri}(R, \hat{\gamma})))$$

4.(d)(i) Compute Q using the routine from 3.b, obtain $\hat{\gamma} = Q^T y$ and compare it to the estimate from `coef(lm(dist ~ Q - 1))`.

```
data(cars)
y <- as.vector(cars$dist)
Q_cars <- q_fun(as.vector(cars$speed), 3)[[1]]
gamma <- crossprod(Q_cars, y) #estimate of gamma as crossprod of Q and y
gamma
```

```
##           [,1]
## [1,] 303.91449
## [2,] 145.55226
## [3,]  22.99576
## [4,]  13.79688
```

```
gamma1 <- coef(lm(cars$dist ~ Q_cars - 1))
gamma1
```

```
##   Q_cars1  Q_cars2  Q_cars3  Q_cars4
## 303.91449 145.55226  22.99576  13.79688
```

4.(d)(ii) Compute $\hat{\beta}$ according to (c) and compare it to the estimate from `coef(lm(dist ~ vandermonde(speed, 3) - 1))`

```
data(cars)
V <- (vander_monde(cars$speed, 3))
colnames(V) <- c("cars1", "cars2", "cars3", "cars4")
```

```
Q <- q_fun(cars$speed, 3)
q_cars <- Q[[1]]
G <- crossprod(q_cars, cars$dist)
G
```

```
##           [,1]
## [1,] 303.91449
## [2,] 145.55226
## [3,]  22.99576
## [4,]  13.79688
```

```
coef(lm(cars$dist ~ q_cars))
```

```
## (Intercept)    q_cars1    q_cars2    q_cars3    q_cars4
##    42.98000         NA    145.55226    22.99576    13.79688
```

```
betal <- coef(lm(cars$dist ~ V - 1))
betal
```

```
##      Vcars1      Vcars2      Vcars3      Vcars4
## -19.50504910  6.80110597 -0.34965781  0.01025205
```

```
qr <- qr(vander_monde(cars$speed, 3))
R <- qr.R(qr)
R
```

```
##           [,1]      [,2]      [,3]      [,4]
## [1,] -7.071068 -108.89444 -1870.7217 -34660.960
## [2,]  0.000000  37.01351  1117.9377  27771.535
## [3,]  0.000000   0.00000  -230.0513 -10089.202
## [4,]  0.000000   0.00000   0.0000  -1345.769
```

```
In <- diag( rep(1, ncol(R)) )
crossprod(inv.upper.tri( R, In , TRUE), G)
```

```
##           [,1]
## [1,] -101.61314068
## [2,]   1.06371154
## [3,]   0.34965781
## [4,]  -0.01025205
```